



Introduction to Deep Learning (CNN)

Prof. Peerapon Vateekul, Ph.D.

Department of Computer Engineering,
Faculty of Engineering, Chulalongkorn University

Peerapon.v@chula.ac.th

www.cp.eng.chula.ac.th/~peerapon/



Outline

- Introduction to Deep Learning
- Convolutional Neural Networks (CNN) [Imaging Task]
- Image Classification
- Object Detection
- Semantic Segmentation







Introduction to Deep Learning

AI = Automation

- 1) Rule-based AI (Symbolic AI)
- 2) Machine Learning




Since an early flush of optimism in the 1950's, smaller subsets of artificial intelligence - first machine learning, then deep learning, a subset of machine learning - have created ever larger disruptions.



<https://mc.ai/machine-learning-basics-artificial-intelligence-machine-learning-and-deep-learning/>

+ Machine Learning





Task1: Supervised learning

Handcrafted features

Training Data



inputs				target
Age	Gender	BodyTemp	Cough	Corona
12	Female	37	Yes	Yes
35	Female	39	No	Yes
32	Male	38	Yes	No

Testing Data



Age	Gender	BodyTemp	Cough	Corona
25	Male	40	No	?

Application: Corona Prediction



Prediction algorithms

- Decision Tree
- (Logistic) Regression
- kNN
- Support Vector Machine
- Neural Networks (NN)
- Deep Learning

BASIC REGRESSION

LINEAR `linear_model.LinearRegression()`
Lots of numerical data     

LOGISTIC `linear_model.LogisticRegression()`
Target variable is categorical  or 

CLASSIFICATION

NEURAL NET `neural_network.MLPClassifier()`
Complex relationships. Prone to overfitting
Basically magic. 

K-NN `neighbors.KNeighborsClassifier()`
Group membership based on proximity 


DECISION TREE `tree.DecisionTreeClassifier()`
If/then/else. Non-contiguous data
Can also be regression  

RANDOM FOREST `ensemble.RandomForestClassifier()`
Find best split randomly
Can also be regression    

SVM `svm.SVC()` `svm.LinearSVC()`
Maximum margin classifier. Fundamental
Data Science algorithm 

NAIVE BAYES `GaussianNB()` `MultinomialNB()` `BernoulliNB()`
Updating knowledge step by step with new info 

Regression – Linear Relationship



weight, coefficient

$$\hat{y} = \hat{w}_0 + \hat{w}_1 x_1 + \hat{w}_2 x_2$$

target intercept input

- The least square method aims to minimize the following term

$$\sum_{\text{training data}} (y_i - \hat{y}_i)^2$$



Logistic Regression (cont.)

Linear Relationship

9

Training Data



inputs				target
Age	Gender	BodyTemp	Cough	Corona
12	Female (0)	37	Yes (1)	Yes
35	Female (0)	39	No (0)	Yes
32	Male (1)	38	Yes (1)	No

$$\text{Logit_score} = w_0 + w_1 * \text{Age} + w_2 * \text{Gender} + w_3 * \text{Temp} + w_4 * \text{Cough}$$


$$\text{Logit_score} = 0.01 - 0.3 * \text{Age} + 0.2 * \text{Gender} + 0.2 * \text{Temp} + 0.9 * \text{Cough}$$

Example

$$\text{Logit_score} = 0.01 - 0.3 * 12 + 0.2 * 0 + 0.2 * 37 + 0.9 * 1 = 4.71$$

prob = 0.9911 → "Yes"


Application: Corona Prediction



$$\hat{p} = \frac{1}{1 + e^{-\text{logit}(\hat{p})}}$$

Neural Networks (universal approximator)

Non-linear relationship




$$\log \left(\frac{\hat{p}}{1-\hat{p}} \right) = \hat{w}_0 + \hat{w}_1 H_1 + \hat{w}_2 H_2 + \hat{w}_3 H_3$$

$$H_1 = \tanh(\hat{w}_{10} + \hat{w}_{11}x_1 + \hat{w}_{12}x_2)$$

$$H_2 = \tanh(\hat{w}_{20} + \hat{w}_{21}x_1 + \hat{w}_{22}x_2)$$

$$H_3 = \tanh(\hat{w}_{30} + \hat{w}_{31}x_1 + \hat{w}_{32}x_2)$$



Stop when?

- Converge (no change in loss)
- Max epochs

Important Params:


- #hidden units, #hidden layers
- Learning rate, Momentum, decay
- Seed number, etc.

Neural Networks (cont.): Training

Non-linear relationship


11

Forward pass (predict)



Backward pass (update weights)

Age	Income	Gender	Province	Corona
25	25,000	Female	Bangkok	Yes
35	50,000	Female	Nontaburi	Yes
32	35,000	Male	Bangkok	No




$$\log \left(\frac{\hat{p}}{1-\hat{p}} \right) = \hat{w}_0 + \hat{w}_1 H_1 + \hat{w}_2 H_2 + \hat{w}_3 H_3$$

$$H_1 = \tanh(\hat{w}_{10} + \hat{w}_{11}x_1 + \hat{w}_{12}x_2)$$

$$H_2 = \tanh(\hat{w}_{20} + \hat{w}_{21}x_1 + \hat{w}_{22}x_2)$$

$$H_3 = \tanh(\hat{w}_{30} + \hat{w}_{31}x_1 + \hat{w}_{32}x_2)$$

Nobel Prize to Backpropagation's Inventors



8 October 2024


The Royal Swedish Academy of Sciences has decided to award the Nobel Prize in Physics 2024 to

John J. Hopfield
Princeton University, NJ, USA

Geoffrey E. Hinton
University of Toronto, Canada

"for foundational discoveries and inventions that enable machine learning with artificial neural networks"

They trained artificial neural networks using physics




John Hopfield invented a network that uses a method for saving and recreating patterns. We can imagine the nodes as pixels. [The Hopfield network](#) utilises physics that describes a material's characteristics due to its atomic spin – a property that makes each atom a tiny magnet.

Geoffrey Hinton used the Hopfield network as the foundation for a new network that uses a different method: [the Boltzmann machine](#). This can learn to recognise characteristic elements in a given type of data.



Handcrafted features

Age	Income	Gender	Province	Corona
25	25,000	Female	Bangkok	Yes
35	50,000	Female	Nontaburi	Yes
32	35,000	Male	Bangkok	No



13

Can we still tell the features (columns)?



shutterstock.com · 451802557




What is Deep Learning (DL)?




Part of the **machine learning** field of learning representations of data. Exceptional effective at learning patterns.




Utilizes learning algorithms that derive meaning out of data by using a **hierarchy** of multiple layers that **mimic the neural networks of our brain**.



Machine Learning



Deep Learning






Deep Learning – Basics (cont.)

What did it learn?


A deep neural network consists of a **hierarchy of layers**, whereby each layer **transforms the input data** into more abstract representations (e.g., edge -> nose -> face). The output layer combines those features to make predictions.



15




Age	Income	Gender	Province	Corona
25	25,000	Female	Bangkok	Yes



x1	x2	x3	x4	Corona
0.7	0.2	-0.5	-0.1	Yes



Deep Learning Application




Speech
Recognition



Computer
Vision



Natural Language
Processing



Type of image tasks

Semantic Segmentation



GRASS, CAT,
TREE, SKY

No objects, just pixels

Classification + Localization



CAT

Single Object

Object Detection



DOG, DOG, CAT

Multiple Object

Instance Segmentation






DOG, DOG, CAT

This image is CC0 public domain

Face Recognition



Smart Medical Diagnosis






Convolutional Neural Networks (CNN)



www.image-net.org

The **ImageNet** project is a large visual database designed for use in visual object recognition software research. Over **14M** URLs of images have been hand-annotated by ImageNet to indicate what objects are pictured on **22K** categories.




ILSVRC

The Image Classification Challenge:


1,000 object classes

1,431,167 images

ImageNet 2017 is the last challenge.




ImageNet Large Scale Visual Recognition Challenge (ILSVRC winners): From SVM to Deep Learning



Wider and deeper networks (Beyond Human)

- ImageNet task




Based on the slide of Aj.Ekapol

Net Large Scale Visual Recognition Challenge, 2014 <https://arxiv.org/abs/1409.0575>

Why now

- Neural Networks has been around since 1990s
- **Big data** – DNN can take advantage of large amounts of data better than other models
- **GPU** – Enable training bigger models possible
- **Deep** – Easier to avoid bad local minima when the model is large



Based on the slide of Aj.Ekapol [/www.kdnuggets.com/2017/06/practical-guide-machine-learning-understand-differentiate-apply.html](http://www.kdnuggets.com/2017/06/practical-guide-machine-learning-understand-differentiate-apply.html)

Application 1: Basic Image Processing

<https://softwaredevelopmentperestroika.wordpress.com/2014/02/11/image-processing-with-python-numpy-scipy-image-convolution/>

การแสดงรูปภาพใน python ด้วย matplotlib

```
import matplotlib.pyplot as plt
import matplotlib.image as mpimg
```

เพื่อความง่ายโปรแกรมที่จะเขียน
จากนี้ไปใช้กับแฟ้ม png เท่านั้น


```
image = mpimg.imread('monument.png')
plt.imshow(image)

plt.show()
```

image เป็นอาร์เรย์ที่แต่ละช่อง
มีค่า 0 ถึง 1 แทนความเข้มของสี


1 pixel = 1 จุดสี มี 3 สีย่อย

0.98762 0.72549 0.35294



It is an array of R,G,B (3 channels)

```
>>> image.shape  
(350, 449, 3)
```




Red
`image[:, :, 0]`

Green
`image[:, :, 1]`

Blue
`image[:, :, 2]`

Basic Image Processing

- `image = mpimg.imread('monument.png')`
- `image.shape` $\rightarrow$ (350, 449, 3)
- `image = 1 - image` (broadcast & element-wise op.)



How's it changed?

Basic Image Processing (1): Negative Image


```
import matplotlib.pyplot as plt
import matplotlib.image as mpimg

image = mpimg.imread('monument.png')
plt.subplot(1, 2, 1)
plt.imshow(image)
```

```
negative = 1 - image
plt.subplot(1, 2, 2)
plt.imshow(negative)
```

```
plt.show()
```

$1 - \text{image}$ = Invert Tone
(White $\rightarrow$ Black, Yellow $\rightarrow$ Blue, ...)
broadcast & element-wise subtract




Basic Image Processing (2): Grayscale Image


```
import numpy as np
import matplotlib.pyplot as plt
import matplotlib.image as mpimg

image = mpimg.imread('monument.png')
plt.subplot(1, 2, 1)
plt.imshow(image)
gray = np.ndarray(image.shape)
gray[:, :, 0] = \
gray[:, :, 1] = \
gray[:, :, 2] = (image[:, :, 0]+image[:, :, 1]+image[:, :, 2]) / 3
plt.subplot(1, 2, 2)
plt.imshow(gray)


plt.show()
```



Another Grayscale Image: $0.299*R + 0.587*G + 0.114*B$




Convolution



Changing the kernel (filter) results in different image processing outcomes


Image Convolution







Convolutional Neural Networks (CNN)

34





CNN (cont.)






Convolutional Neural Networks (CNN)

1) Convolutional Layer

- 1 • Convolutional Layer
- 2 • Pooling Layer
- 3 • Fully-Connected Layer





Layer 1: Convolutional Layer

1 _{x1}	1 _{x0}	1 _{x1}	0	0
0 _{x0}	1 _{x1}	1 _{x0}	1	0
0 _{x1}	0 _{x0}	1 _{x1}	1	1
0	0	1	1	0
0	1	1	0	0


Image

(3*3 filter)

1	0	1
0	1	0
1	0	1

4		

Convolved
Feature





com

Input Volume (+pad 1) (7x7x3)

$x[:, :, 0]$

0	0	0	0	0	0	0
0	1	0	2	2	2	0
0	0	0	0	0	0	0
0	2	0	2	2	2	0
0	1	0	0	0	0	0
0	1	0	0	2	1	0
0	0	0	0	0	0	0

$x[:, :, 1]$

0	0	0	0	0	0	0
0	1	0	1	0	1	0
0	0	0	0	0	1	0
0	0	0	2	0	0	0
0	2	0	0	0	0	0
0	0	0	1	0	0	0
0	0	0	0	0	0	0

$x[:, :, 2]$

0	0	0	0	0	0	0
0	0	2	2	0	0	0
0	0	0	0	0	1	0
0	1	1	2	0	2	0
0	2	0	0	0	0	0
0	0	1	1	0	1	0
0	0	0	0	0	0	0

Filter W0 (3x3x3)

$w0[:, :, 0]$

1	0	0
-1	0	0
0	-1	1

$w0[:, :, 1]$

-1	-1	0
-1	-1	1
1	0	0

$w0[:, :, 2]$

1	1	0
0	-1	1
1	1	1

Bias b0 (1x1x1)

$b0[:, :, 0]$

1

Filter W1 (3x3x3)

$w1[:, :, 0]$

1	-1	0
-1	0	0
0	1	-1

$w1[:, :, 1]$

0	1	-1
0	0	0
0	1	1

$w1[:, :, 2]$

0	-1	0
1	1	-1
-1	0	1

Bias b1 (1x1x1)

$b1[:, :, 0]$

0

Output Volume (3x3x2)

$o[:, :, 0]$

2	-2	-1
2	-3	-3
2	-1	-2

$o[:, :, 1]$

-2	4	-1
3	3	0
-2	2	-1

$$\begin{aligned}
 o(2,2,0) &= \sum x[:, :, 0] \times w[:, :, 0] + \sum x[:, :, 1] \times w[:, :, 1] + \sum x[:, :, 2] \times w[:, :, 2] + b_0 \\
 &= 0 \times 1 + 0 \times 0 + 0 \times 0 + 2 \times (-1) + 1 \times 0 + 0 \times 0 + 0 \times 0 + 0 \times (-1) + 0 \times 1 \\
 &\quad + 0 \times (-1) + 0 \times (-1) + 0 \times 0 + 0 \times (-1) + 0 \times (-1) + 0 \times 1 + 0 \times 1 + 0 \times 0 + 0 \times 1 \\
 &\quad + 0 \times 1 + 0 \times 1 + 0 \times 0 + 0 \times 0 + 1 \times (-1) + 0 \times 1 + 0 \times 1 + 0 \times 1 + 0 \times 1 \\
 &\quad + 1 \\
 &= -2
 \end{aligned}$$

CNNs: 1) Convolutional Layer

Feature Extraction

<http://cs231n.github.io/convolutional-networks/>

(3*3 filter)

$$w^T x + b$$

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Image

1	0	1
0	1	0
1	0	1

4		

Convolved
Feature


(feature map)

Filter size=(3*3), #Filters=1, Stride=1, Padding=0
Output size = (5 + 2(0) - 3) / 1 + 1 = 3

Summary: **4 parameters** for convolutional layer

- (1) Filter size, (2) #Filters, (3) Stride, (4) Padding


$$\text{Output size} = (N + 2 * P - F) / S + 1$$





CNNs: 1) Convolutional Layer (cont.)

Non-Linear Activation Function






Convolutional Neural Networks (CNN)

2) Pooling Layer

- 1 • Convolutional Layer
- 2 • Pooling Layer
- 3 • Fully-Connected Layer




CNNs: 2) Pooling Layer

<http://cs231n.github.io/convolutional-networks/>


Summary: 2 parameters for pooling layer

- Filter size, Stride

■ Feature selection (reduction); downsampling



two more layers to go: POOL/FC







Convolutional Neural Networks

Embedding Vector

x1	x2	x3	x4	Corona
0.7	0.2	-0.5	-0.1	Yes



- 1 • Convolutional Layer
- 2 • Pooling Layer
- 3 • Fully-Connected Layer







Convolutional Neural Networks (CNN)

3) Fully-Connected Layer (FC) = Neural Networks

- 1 • Convolutional Layer
- 2 • Pooling Layer
- 3 • Fully-Connected Layer



ImageNet Large Scale Visual Recognition Challenge (ILSVRC 2nd runner-up in 2014): VGG19




CNN Example: VGGNet (2014)

[Simonyan and Zisserman, 2014]


Small filters, Deeper networks

8 layers (AlexNet) → 16-19 layers (VGG16Net)

Only 3x3 CONV Stride 1, pad 1
And 2x2 MAX POOL stride 2



AlexNet



VGG16




VGG19

Network	A	B	C	D	E
Number of parameters (Millions)	133	133	134	138	144

47

ImageNet Large Scale Visual Recognition Challenge (ILSVRC 2nd runner-up in 2014): Deeper than VGG19



Install

Learn

API

Resources

Community

Why TensorFlow

Overview

Input

Model

Sequential

activations

applications

Overview

DenseNet121

DenseNet169

DenseNet201

EfficientNetB0

EfficientNetB1

EfficientNetB2

EfficientNetB3

EfficientNetB4

EfficientNetB5

EfficientNetB6

EfficientNetB7

InceptionResNetV2

InceptionV3

MobileNet

MobileNetV2

MobileNetV3Large

MobileNetV3Small

NASNetLarge

NASNetMobile

Modules

densenet module: DenseNet models for Keras.

efficientnet module: EfficientNet models for Keras.

imagenet_utils module: Utilities for ImageNet data preprocessing & prediction decoding.

inception_resnet_v2 module: Inception-ResNet V2 model for Keras.

inception_v3 module: Inception V3 model for Keras.

mobilenet module: MobileNet v1 models for Keras.

mobilenet_v2 module: MobileNet v2 models for Keras.

nasnet module: NASNet-A models for Keras.

resnet module: ResNet models for Keras.

resnet50 module: Public API for tf.keras.applications.resnet50 namespace.

resnet_v2 module: ResNet v2 models for Keras.

vgg16 module: VGG16 model for Keras.

vgg19 module: VGG19 model for Keras.

Model 1

- VGG19 (random initialized weights) + 2 Dense layers + Output layer

```

1 base_model = VGG19(weights=None, include_top=False, input_shape=(224, 224, 3))
2
3 for layer in base_model.layers:
4     layer.trainable = True
5
6 x = base_model.output
7 x = Flatten()(x)
8 x = Dense(1024)(x)
9 x = Dropout(0.5)(x)
10 x = Dense(512)(x)
11 x = Dropout(0.5)(x)
12 output = Dense(num_class, activation='softmax')(x)
13

```

Package Reference

torchvision.datasets

torchvision.io

torchvision.models

torchvision.ops

torchvision.transforms

torchvision.utils

PyTorch Libraries

PyTorch

torchaudio

torchtext

torchvision

TorchElastic

TorchServe

PyTorch on XLA Devices

TORCHVISION.MODELS

The models subpackage contains definitions of models for addressing different tasks, including: image classification, pixelwise semantic segmentation, object detection, instance segmentation, person keypoint detection and video classification.

Classification

The models subpackage contains definitions for the following model architectures for image classification:

- AlexNet
- VGG
- ResNet
- SqueezeNet
- DenseNet
- Inception v3
- GoogLeNet
- ShuffleNet v2
- MobileNetV2
- MobileNetV3
- ResNeXt
- Wide ResNet
- MNASNet

We provide pre-trained models, using the PyTorch `torch.utils.model_zoo`. These can be constructed by passing `pretrained=True`:

```
import torchvision.models as models
resnet18 = models.resnet18(pretrained=True)
alexnet = models.alexnet(pretrained=True)
squeezenet = models.squeezenet1_0(pretrained=True)
vgg16 = models.vgg16(pretrained=True)
densenet = models.densenet161(pretrained=True)
inception = models.inception_v3(pretrained=True)
googlenet = models.googlenet(pretrained=True)
shufflenet = models.shufflenet_v2_x1_0(pretrained=True)
mobilenet_v2 = models.mobilenet_v2(pretrained=True)
mobilenet_v3_large = models.mobilenet_v3_large(pretrained=True)
mobilenet_v3_small = models.mobilenet_v3_small(pretrained=True)
resnext50_32x4d = models.resnext50_32x4d(pretrained=True)
wide_resnet50_2 = models.wide_resnet50_2(pretrained=True)
mnasnet = models.mnasnet1_0(pretrained=True)
```



v1.1.0 ▼

🏠 pytorch-transformers

🌟 Star 131,473

Search docs

NOTES

Installation

Quickstart

Pretrained models

Examples

Notebooks

Loading Google AI or OpenAI pre-trained weights or PyTorch dump

Serialization best-practices

Converting Tensorflow Checkpoints

Migrating from pytorch-pretrained-bert

BERTology

ⓘ You are viewing legacy docs. Go to [latest documentation](#) instead.

Docs » Pytorch-Transformers

Pytorch-Transformers

PyTorch-Transformers is a library of state-of-the-art pre-trained models for Natural Language Processing (NLP).

The library currently contains PyTorch implementations, pre-trained model weights, usage scripts and conversion utilities for the following models:

1. [BERT](#) (from Google) released with the paper [BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding](#) by Jacob Devlin, Ming-Wei Lee and Kristina Toutanova.
2. [GPT](#) (from OpenAI) released with the paper [Improving Language Understanding by Generative Pre-Training](#) by Alec Radford, Karthik Narasimhan, Tim Salim Sutskever.
3. [GPT-2](#) (from OpenAI) released with the paper [Language Models are Unsupervised Multitask Learners](#) by Alec Radford*, Jeffrey Wu*, Rewon Child, David Lu and Ilya Sutskever**.
4. [Transformer-XL](#) (from Google/CMU) released with the paper [Transformer-XL: Attentive Language Models Beyond a Fixed-Length Context](#) by Zihang Dai*, Z Yang, Jaime Carbonell, Quoc V. Le, Ruslan Salakhutdinov.
5. [XLNet](#) (from Google/CMU) released with the paper [XLNet: Generalized Autoregressive Pretraining for Language Understanding](#) by Zhilin Yang*, Zihang Dai, Jaime Carbonell, Ruslan Salakhutdinov, Quoc V. Le.
6. [XLM](#) (from Facebook) released together with the paper [Cross-lingual Language Model Pretraining](#) by Guillaume Lample and Alexis Conneau.

Notes

- [Installation](#)

<https://www.martincid.com/technology-sv/nvidia-reportedly-buys-hugging-face-12-9-billion/>



AI

Nvidia bets \$12.9B that owning Hugging Face won't kill open-source AI



Nvidia. Photo: 4300streetcar / CC BY 4.0, via Wikimedia Commons



By **Adrian Kessler**

31 August 2026, 3:37 am · 3 min read



Image Classification Tasks

Classification




CAT

**Classification
+ Localization**



CAT

Object Detection



CAT, DOG, DUCK

**Instance
Segmentation**



CAT, DOG, DUCK


Single object

Multiple objects




Image Classification (recap)

Convolutional Neural Network






SOTA of Image Classification



https://blog.csdn.net/qg_34216467/article/details/83061692



<https://theaisummer.com/cnn-architectures/>




Common Image Classification Models

56

Model	Year	Team / Organization
ResNet	2015	Microsoft Research
MobileNet	2017	Google
EfficientNet	2019	Google
Vision Transformer (ViT)	2020	Google Research
Swin Transformer	2021	Microsoft Research
ConvNeXt	2022	Meta AI

Image Classification on ImageNet



<https://paperswithcode.com/sota/image-classification-on-imagenet>

+ Image classification dashboard (2026)

1. **CNN / ViT** → **Train/Fine-tune**
2. **DINOv3** → **Freeze backbone + Train head**
3. **CLIP / SigLIP** → **Zero-shot**

■ CNN → Vision Transformer (ViT) → Foundation Models / VLMs

- 2021–2022 → CNN → ViT: CNNs (e.g., EfficientNetV2) remained strong, while Vision Transformers increasingly dominated at scale.
- 2022–2024 → Large ViTs + Vision–Language pretraining: CLIP & CoCa enabled powerful transferable representations.
- 2025 → Foundation Vision Models: DINOv3 scaled self-supervised learning to provide strong general-purpose visual representations.
- 2025–2026 → Foundation Models + VLMs: Vision-only models (DINOv3) and Vision–Language models (SigLIP 2, CLIP-family) coexist as major pretraining paradigms.

■ 2026 Key Takeaway:

- Retrain a foundation encoder → adapt / fine-tune, rather than train an image classifier from scratch.
- CNN / ConvNeXt → Strong inductive bias; efficient and still relevant.
- ViT / DINOv3 → Vision-only (image); self-supervised foundation encoder with transferable features.
- CLIP / SigLIP 2 / CoCa → Vision–Language (image + text); supports zero-shot classification, retrieval, and multimodal tasks.



EfficientNet (May, 2019) → EffNetV2 (2021)

58

EfficientNet: Rethinking Model Scaling for Convolutional Neural Networks

Mingxing Tan¹ Quoc V. Le¹

Abstract

Convolutional Neural Networks (ConvNets) are commonly developed at a fixed resource budget, and then scaled up for better accuracy if more resources are available. In this paper, we systematically study model scaling and identify that carefully balancing network depth, width, and resolution can lead to better performance. Based on this observation, we propose a new scaling method that uniformly scales all dimensions of depth/width/resolution using a simple yet highly effective *compound coefficient*. We demonstrate the effectiveness of this method on scaling up MobileNets and ResNet.

To go even further, we use neural architecture





Figure 1. Model Size vs. ImageNet Accuracy. All numbers are



(a) Training efficiency.

	EfficientNet (2019)	ResNet-RS (2021)	DeiT/ViT (2021)	EfficientNetV2 (ours)
Top-1 Acc.	84.3%	84.0%	83.1%	83.9%
Parameters	43M	164M	86M	24M



(b) Parameter efficiency.

<http://proceedings.mlr.press/v97/tan19a/tan19a.pdf>

<https://twitter.com/TensorFlow/status/1423359562724175873/photo/1>

+ Vision Transformer (ViT) (ICLR2021)


🤗 Hugging Face





<https://medium.com/machine-intelligence-and-deep-learning-lab/vit-vision-transformer-cc56c8071a20>

+ ConvNeXt (CVPR2022) → ConvNeXt V2 (2023)

60




Swin Transformer Block



ConvNeXt V1 Block

ConvNeXt V2 Block



<https://github.com/facebookresearch/ConvNeXt-V2>

<https://github.com/facebookresearch/ConvNeXt>



Open Source

DINOv3: Self-supervised learning for vision at unprecedented scale

August 14, 2025 • ⌚ 11 minute read

61

- DINOv3: Self-Supervised Vision Transformer (Meta AI, 2025)
- Vision-only (**encoder-only backbone**) (not multimodal, unlike CLIP/CoCa).
- Self-supervised learning → no human labels needed.
- **Massive scale: up to 7B parameters, trained on 1.7B images.**
- Key innovations:
 - Gram Anchoring → stabilizes feature learning.
 - Axial RoPE → improved positional encoding.
 - Resolution robustness → better handling of varied image sizes.
- Strengths:
 - Provides general-purpose visual representations.
 - **Excels in downstream tasks (classification, detection, segmentation).**
 - Foundation model role for vision, similar to LLMs in NLP.


[Link](#)




SOTA of Object Detection

62

One and two stage detectors



V7 Labs




https://medium.com/@jonathan_hui/real-time-object-detection-with-yolo-yolov2-28b1b93e2088




Object Detection

One and two stage detectors



V7



63

		units	
		absolute	normalized
corner coordinate	(left, top, right, bottom)	Pascal VOC	Albumentations
	(left, top, width, height)	COCO	
	(center_x, center_y, width, height)	CreateML	YOLO

<https://dragoneye.ai/blog/a-guide-to-bounding-box-formats/>

+ YOLO Open Images in New York

YOLO Open Images in New York



Joseph Redmon

6.95K subscribers

Subscribe

13,911 views Nov 13, 2017

No description has been added to this video.

64



+ YOLO v5

Example of YOLO v5 detection on video file




Luiz doleron
127 subscribers

Subscribe

5,781 views Jan 23, 2022

Video example of YOLOv5 performing object detection.

- Check out the explanation at [🔗 / detecting-objects-with-yolov5-opencv-pytho...](#)
- Github repository is here: <https://github.com/doleron/yolov5-ope...>






What is YOLO?

You Only Look Once (YOLO) is a one-stage, end-to-end object detection framework that predicts bounding boxes and class probabilities in a single forward pass of a neural network.

Unlike earlier object detection approaches that adapted image classifiers into multi-step detection pipelines, YOLO treats object detection as a single regression problem.



News

- We are pleased to announce the [LVIS 2021 Challenge and Workshop](#) to be held at ICCV.
- Please note that there will not be a COCO 2021 Challenge, instead, we encourage people to participate in the LVIS 2021 Challenge.
- We have partnered with the team behind the open-source tool [FiftyOne](#) to make it easier to download, visualize, and evaluate COCO
- [FiftyOne](#) is an open-source tool facilitating visualization and access to COCO data resources and serves as an evaluation tool for model analysis on COCO.

What is COCO?



COCO is a large-scale object detection, segmentation, and captioning dataset. COCO has several features:

- ✓ **Object segmentation**
- ✓ **Recognition in context**
- ✓ **Superpixel stuff segmentation**
- ✓ **330K images (>200K labeled)**
- ✓ **1.5 million object instances**
- ✓ **80 object categories**
- ✓ **91 stuff categories**
- ✓ **5 captions per image**
- ✓ **250,000 people with keypoints**

<https://cocodataset.org/#home>

Collaborators

Tsung-Yi Lin Google Brain
Genevieve Patterson MSR, Trash TV
Matteo R. Ronchi Caltech
Yin Cui Google
Michael Maire TTI-Chicago
Serge Belongie Cornell Tech
Lubomir Bourdev WaveOne, Inc.
Ross Girshick FAIR
James Hays Georgia Tech
Pietro Perona Caltech
Deva Ramanan CMU
Larry Zitnick FAIR
Piotr Dollár FAIR

Sponsors




**CVDF****Microsoft****ultralytics/cfg/datasets/coco8.yaml**

```
# Classes
names:
  0: person
  1: bicycle
  2: car
  3: motorcycle
  4: airplane
  5: bus
  6: train
  7: truck
  8: boat
  9: traffic light
  10: fire hydrant
  11: stop sign
  12: parking meter
  13: bench
  14: bird
  15: cat
  16: dog
  17: horse
  18: sheep
  19: cow
  20: elephant
  21: bear
  22: zebra
  23: giraffe
```

79 classes

<https://docs.ultralytics.com/datasets/detect/#ultralytics-yolo-format>

A photograph showing a person from behind, wearing a bright blue jacket, a dark hat, and dark pants, walking through a dry, grassy field. Two dogs are walking alongside them: a medium-sized, scruffy dog with grey and black fur, and a dark-colored horse. In the background, there are patches of snow on the ground and a simple black metal frame structure.




<https://blog.paperspace.com/how-to-implement-a-yolo-object-detector-in-pytorch/>

The Evolution of YOLO: 2016 - 2025

YOLO Evolution: 2016–2025

Four paradigm phases culminating in 2025 bifurcation toward accuracy and deployment



Ultralytics YOLO → production-ready, practical framework (YOLOv5 → YOLOv8 → YOLO11 → YOLO26).
Academic YOLO → new research architectures (YOLOv9, v10, v11, v12).


+ Ultralytic's YOLO

YOLOv5 (2020) → YOLOv8 (2023) → YOLO 11 (2024) → YOLO 26 (2025)


70



DOWNLOAD THE APP





English | 简体中文



Ultralytics YOLOv8 is a cutting-edge, state-of-the-art (SOTA) model that builds upon the success of previous YOLO versions and introduces new features and improvements to further boost performance and flexibility. YOLOv8 is designed to be fast, accurate, and easy to use, making it an excellent choice for a wide range of object detection and tracking, instance segmentation, image classification and pose estimation tasks.


We hope that the resources here will help you get the most out of YOLOv8. Please browse the YOLOv8 [Docs](#) for details, raise an issue on [GitHub](#) for support, and join our [Discord](#) community for questions and discussions!

To request an Enterprise License please complete the form at [Ultralytics Licensing](#).



Models

YOLOv8 [Detect](#), [Segment](#) and [Pose](#) models pretrained on the [COCO](#) dataset are available here, as well as YOLOv8 [Classify](#) models pretrained on the [ImageNet](#) dataset. [Track](#) mode is available for all Detect, Segment and Pose models.



All [Models](#) download automatically from the latest Ultralytics [release](#) on first use.

+ YOLO Models (2024–2026) Timeline

71

Model	Release	Team / University / Organization	Key Highlight
YOLO26	Jan 2026	Ultralytics	Native end-to-end, NMS-free ; DFL-free head; MuSGD; optimized for edge deployment
YOLO12	Feb 2025	Yunjie Tian, Qixiang Ye (UCAS), David Doermann (University at Buffalo)	Attention-centric YOLO ; Area Attention + R-ELAN
YOLO11	Sep 2024	Ultralytics	Improved accuracy–efficiency; detection, segmentation, pose, classification & OBB
YOLOv10	May 2024	Tsinghua University	End-to-end, NMS-free detection; consistent dual assignments
YOLOv9	Feb 2024	Institute of Information Science, Academia Sinica, Taiwan	PGI + GELAN for better gradient flow and efficiency

Evolution: Better gradients → End-to-end / NMS-free → Multi-task efficiency → Attention → Edge-first deployment


+ DETR (Detection Transformer) [Meta, 2020]

72

DETR: End-to-End Object Detection with Transformers

Support Ukraine

PyTorch training code and pretrained models for **DETR** (**DE**tection **TR**ansformer). We replace the full complex hand-crafted object detection pipeline with a Transformer, and match Faster R-CNN with a ResNet-50, obtaining **42 AP** on COCO using half the computation power (FLOPs) and the same number of parameters. Inference in 50 lines of PyTorch.



- 1. Deformable DETR (2020)
■ Team: Microsoft Research Asia
- 2. Conditional DETR (2021)
■ Team: Chinese Academy of Sciences
- 3. Anchor DETR (2021)
■ Team: SenseTime and Shanghai Jiao Tong University
- 4. Efficient DETR (2021)
■ Team: UC Berkeley and Google Research
- 5. DINO (DETR with Improved Non-Autoregressive Decoding) (2022)
■ Team: Meta AI (formerly Facebook AI Research)
- 6. RT-DETR (Real-Time DETR) (2024)
■ Team: Alibaba DAMO Academy



Object Detection Landscape (2026)

73

■ NN-based → Transformer-based → End-to-End Detection

- 2016–2020 → CNN-based detectors: Faster R-CNN, SSD, YOLO established two-stage vs. one-stage detection.
- 2020–2022 → Transformer detectors: [DETR](#) / Deformable DETR introduced end-to-end set prediction without anchors or traditional NMS.
- 2023–2024 → Real-time Transformers: [RT-DETR](#) brought Transformer-based detection to real-time applications.
- 2024–2025 → End-to-End YOLO: YOLOv10 / YOLO11 / YOLO12 improved efficiency, attention, and end-to-end detection.
- 2026 → Efficient End-to-End Detection: [YOLO26](#) emphasizes NMS-free inference and edge deployment.

■ 2026 Key Takeaway

- Modern object detection is shifting toward end-to-end prediction: Image → Boxes + Classes
- Faster R-CNN → Two-stage CNN; region proposals + classification.
- YOLO → One-stage, real-time detector; modern versions increasingly end-to-end.
- DETR / RT-DETR → Transformer-based, end-to-end object detection.
- YOLO26 → End-to-end + NMS-free + edge-optimized real-time detection.

■ Evolution:



- Two-stage CNN → One-stage YOLO → Transformer (DETR) → End-to-End / NMS-free Detection

[Personal open source] - -

Real - time Semantic Segmentation ddrnet

2021-11-09 04:24:01

[Nongfu Mountain Spring 2]



- **Real-time semantic segmentation**
- BiSeNet V1, V2
- DDRNet



<https://chowdera.com/2021/11/20211109042359120o.html>



Semantic Segmentation


75

- Semantic segmentation = pixel-wise classification
- Each pixel gets a class label (road, car, polyp, tumor, sky, background).



UNet: Encoder-Decoder Network: Encoder, Decoder, Skip Connections

U-Net



Deconvolutional layer



$$\mathcal{L} = \lambda_1 \cdot \text{CE} + \lambda_2 \cdot \text{Dice} ; \text{CE (pixel loss) \& Dice (region loss; semantic)}$$

Architecture of the U-net for a given input image. The blue boxes correspond to feature maps blocks with their denoted shapes. The white boxes correspond to the copied and cropped feature maps.

Source: [O. Ronneberger et al. \(2015\)](#)

+ Image Segmentation Landscape (2026)

■ CNN → Transformer → Foundation Models

- 2015–2020 → CNN-based: **FCN, U-Net, DeepLab** → encoder–decoder architectures for semantic segmentation.
- 2021–2022 → Transformer-based: **SegFormer, Swin, Mask2Former** → global attention; semantic, instance & panoptic segmentation.
- 2023 → Foundation Models: **SAM** → introduced promptable segmentation using points, boxes, or masks.
- 2024 → Image + Video: **SAM 2** → extended promptable segmentation to images and videos.
- 2025–2026 → Foundation Vision: **DINOv3 + segmentation head** → strong pretrained features for dense prediction.

■ 2026 Key Takeaway

- Task-specific training → Pretrained foundation models + lightweight adaptation / prompting
- U-Net / DeepLab → CNN-based → Train / Fine-tune
- SegFormer / Mask2Former → Transformer-based → Train / Fine-tune
- DINOv3 + Head → Freeze backbone + Train segmentation head
- SAM / SAM 2 → Promptable segmentation / Zero-shot

■ Evolution:

- CNN → Transformer → Foundation Model → Promptable Segmentation

Introducing SAM 2: The next generation of Meta Segment Anything Model for videos and images

July 29, 2024 • ⌚ 15 minute read

INTRODUCING

Meta Segment Anything Model 2 for videos and images



AI at Meta

<https://ai.meta.com/blog/segment-anything-2/>


+ Ultralytic's YOLO (extra segmentation model)

YOLOv5 (2020) → YOLOv8 (2023) → Ultralytics YOLO 11 (2024)

79





English | 简体中文



Ultralytics YOLOv8 is a cutting-edge, state-of-the-art (SOTA) model that builds upon the success of previous YOLO versions and introduces new features and improvements to further boost performance and flexibility. YOLOv8 is designed to be fast, accurate, and easy to use, making it an excellent choice for a wide range of object detection and tracking, instance segmentation, image classification and pose estimation tasks.


We hope that the resources here will help you get the most out of YOLOv8. Please browse the YOLOv8 [Docs](#) for details, raise an issue on [GitHub](#) for support, and join our [Discord](#) community for questions and discussions!

To request an Enterprise License please complete the form at [Ultralytics Licensing](#).



Models

YOLOv8 [Detect](#), [Segment](#) and [Pose](#) models pretrained on the [COCO](#) dataset are available here, as well as YOLOv8 [Classify](#) models pretrained on the [ImageNet](#) dataset. [Track](#) mode is available for all Detect, Segment and Pose models.



All [Models](#) download automatically from the latest Ultralytics [release](#) on first use.



YOLOv8-Seg: Object Detection → Segment

80








Fig.3 YOLOv8-Seg architecture; a backbone, b neck, c head, d Proto module, e C2f module and f SPPF module

[Link](#)



Real-Time Colonic Polyp Detection

An AI-assisted solution that aims to improve colonic polyp detection in real-time. It is compatible with all standard colonoscope systems.



09/11/2021
13:53:24



COLON

*1/200
Lv+2 AUTO

HT NR

SE

f

*
L

3.8

S1: F+T
12.0 S2: LM
12.0 S3: CAD
S4:

EC-760R-V/L

3C729K035

BL-7000

CHULALONGKORN HOS



Name:


Sex: Age:
D.O.B.:
18/05/2021
08:29:36

■ ■ □ / --- (0/1)
Eh:B8 Cm:1



Comment:

NBI

threshold : 0.4
class : NBI



Cholangioscopy





Appendix